



Deep Learning

Fundamentals

Machine Learning

*Dongsuk Yook
Artificial Intelligence Laboratory
Korea University*

Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications

Class Objectives

- Understanding the fundamentals of deep learning
- Being able to build a basic deep neural network for a given application


Contents

- Perceptrons
 - Neurons
 - Learning for Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications

Neurons

□ Biological neural networks


- Approximately 10^{11} interconnected *neurons*, each connected to 10–100,000 others through *synapses*
- More than 10^{14} synaptic connections
- Highly complex and massively parallel



Perceptrons

□ Perceptron [McCulloch and Pitts 1943, Rosenblatt 1957]

- Input x_j
- Output y
- Connection weight (or synaptic weight) w_j




□ Linear regression

- $y = \sum_j w_j x_j + w_0 = \mathbf{w}^\top \mathbf{x}$

; $\mathbf{x} \equiv [1, x_1, \dots, x_d]$
 w_0 : bias


□ Binary classification

- Linear discriminant function using a *hard threshold* function
- $y = \begin{cases} 1 & \text{if } \mathbf{w}^\top \mathbf{x} > 0 \\ 0 & \text{otherwise} \end{cases}$



□ Sigmoid function for posterior probability

- $y = \frac{1}{1+\exp(-\mathbf{w}^\top \mathbf{x})}$



; $\exp(-\mathbf{w}^\top \mathbf{x}) \equiv e^{-\mathbf{w}^\top \mathbf{x}}$


Multiple Perceptrons

- Multiple perceptrons
 - Input x_j
 - Output y_i
 - Weight w_{ij} (between input x_j and output y_i)

- Multivariate linear regression
 - $y_i = \sum_j w_{ij}x_j + w_{i0} = \mathbf{w}_i^\top \mathbf{x}$
 - $\mathbf{y} = \mathbf{W}\mathbf{x}$

- Multiclass classification
 - $\mathbf{y} = \mathbf{W}\mathbf{x}$
 - $\mathbf{x} \in c_k$ if $k = \arg \max_i y_i$

- Softmax function for posterior probability
 - $s_i = \sum_j w_{ij}x_j + w_{i0} = \mathbf{w}_i^\top \mathbf{x}$
 - $y_i = \frac{\exp(s_i)}{\sum_i \exp(s_i)}$





Contents

- Perceptrons
 - Neurons
 - Learning for Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications

Learning for a Linear Output Node


- Stochastic gradient descent (SGD) for a linear output node
 - $E(\mathbf{w}|\mathbf{x}, r) \equiv \frac{1}{2}(r - y)^2$; mean squared error (MSE)
 - $\frac{\partial E}{\partial w_j} = \frac{\partial E}{\partial y} \frac{\partial y}{\partial w_j}$; $y = \mathbf{w}^\top \mathbf{x} = \sum_j w_j x_j + w_0$
$$= \underbrace{\frac{\partial E}{\partial y}}_{\frac{\partial E}{\partial y}} \underbrace{\frac{\partial y}{\partial w_j}}_{x_j}$$
 - $w_j \leftarrow w_j + \eta(r - y)x_j$; η : learning rate or step size



Learning for Multiple Linear Output Nodes


- Stochastic gradient descent for multiple linear output nodes

- $E(\mathbf{W}|\mathbf{x}, \mathbf{r}) \equiv \sum_k \underbrace{\frac{1}{2}(r_k - y_k)^2}_{E_k}$; $y_i = \mathbf{w}_i^\top \mathbf{x} = \sum_j w_{ij}x_j + w_{i0}$
- $$\begin{aligned}\frac{\partial E}{\partial w_{ij}} &= \sum_k \frac{\partial E_k}{\partial y_i} \frac{\partial y_i}{\partial w_{ij}} \\ &= \underbrace{-(r_i - y_i)}_{\frac{\partial E_i}{\partial y_i}} \underbrace{x_j}_{\frac{\partial y_i}{\partial w_{ij}}}\end{aligned}$$
- $w_{ij} \leftarrow w_{ij} + \eta(r_i - y_i)x_j$



Entropy


- Entropy
 - $\mathbb{H}(p) \equiv -\sum_x p(x) \log_2 p(x)$; $\mathbb{H}(x)$
 - e.g., $-[\theta \log_2 \theta + (1-\theta) \log_2 (1-\theta)]$
- Cross entropy
 - $\mathbb{H}(p, q) \equiv -\sum_x p(x) \log_2 q(x)$; minimum when $p = q$
(\because Jensen's inequality)
- Kullback-Leibler (KL) divergence (also called *relative entropy*)
 - $\mathbb{KL}(p||q) \equiv \mathbb{H}\left(p(x), \frac{q(x)}{p(x)}\right) = -\sum_x p(x) \log_2 \frac{q(x)}{p(x)}$
 $= -\sum_x p(x) \log_2 q(x) + \sum_x p(x) \log_2 p(x) = \mathbb{H}(p, q) - \mathbb{H}(p)$
- Conditional entropy
 - $\mathbb{H}(x|y) \equiv \mathbb{H}(p(x,y), p(x|y)) = -\sum_{x,y} \widehat{p(x,y)} \log_2 p(x|y)$
 $= -\sum_y \sum_x p(y) p(x|y) \log_2 p(x|y) = -\sum_y p(y) \mathbb{H}(p(x|y))$
- Mutual information (MI, also called *information gain*)
 - $\mathbb{I}(x, y) \equiv \mathbb{KL}(p(x,y)||p(x)p(y)) = \mathbb{H}(x) - \mathbb{H}(x|y) = \mathbb{H}(y) - \mathbb{H}(y|x)$



Learning for a Sigmoid Output Node

- Stochastic gradient descent for a sigmoid output node

- $E(\mathbf{w}|\mathbf{x}, r) \equiv -r \log y - (1 - r) \log(1 - y)$; cross entropy
 - $$\begin{aligned}\frac{\partial E}{\partial w_i} &= \frac{\partial E}{\partial y} \frac{\partial y}{\partial s} \frac{\partial s}{\partial w_i} \\ &= \underbrace{\left(-\frac{r}{y} + \frac{1-r}{1-y}\right)}_{\frac{\partial E}{\partial y}} \underbrace{y(1-y)}_{\frac{\partial y}{\partial s}} \underbrace{x_i}_{\frac{\partial s}{\partial w_i}} \\ &= -(r - y)x_i \\ &w_i \leftarrow w_i + \eta(r - y)x_i\end{aligned}$$
- $; y = \frac{1}{1+\exp(-s)}, s = \mathbf{w}^\top \mathbf{x}$
- $$\begin{aligned}y' &= -\frac{1}{(1+\exp(-s))^2} \exp(-s) (-1) \\ &= \frac{1}{1+\exp(-s)} \frac{\exp(-s)}{1+\exp(-s)} \\ &= \frac{1}{1+\exp(-s)} \left(\frac{1+\exp(-s)}{1+\exp(-s)} - \frac{1}{1+\exp(-s)} \right) \\ &= y(1 - y)\end{aligned}$$




Learning for Softmax Output Nodes

- Stochastic gradient descent for a softmax output node

- $$E(\mathbf{W}|\mathbf{x}, \mathbf{r}) \equiv -\sum_k r_k \underbrace{\log y_k}_{E_k}$$

$$\therefore y_k = \frac{\exp(s_k)}{\sum_{\hat{k}} \exp(s_{\hat{k}})}, s_i = \mathbf{w}_i^T \mathbf{x}$$
- $$\begin{aligned} \frac{\partial E}{\partial w_{ij}} &= -\sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial s_i} \frac{\partial s_i}{\partial w_{ij}} \\ &= - \left(\underbrace{\sum_{k \neq i} \frac{r_k}{y_k} \left(-\frac{\exp(s_k) \exp(s_i)}{\left(\sum_{\hat{k}} \exp(s_{\hat{k}}) \right)^2} \right)}_{\frac{\partial E_k}{\partial y_k}} + \underbrace{\frac{r_i}{y_i} \left(\frac{\exp(s_i)}{\sum_i \exp(s_i)} - \frac{\exp(s_i) \exp(s_i)}{\left(\sum_i \exp(s_i) \right)^2} \right)}_{\frac{\partial E_i}{\partial y_i}} \right) \frac{x_j}{\frac{\partial s_i}{\partial w_{ij}}} \\ &= - \left(\sum_{k \neq i} \frac{r_k}{y_k} (-y_k y_i) + \frac{r_i}{y_i} (y_i - y_i^2) \right) x_j \\ &= -(\sum_{k \neq i} r_k (-y_i) + r_i (1 - y_i)) x_j \\ &= -(\sum_k r_k (-y_i) + r_i) x_j \\ &= -(r_i - y_i) x_j \\ \end{aligned}$$
- $w_{ij} \leftarrow w_{ij} + \eta(r_i - y_i)x_j$



Stochastic Gradient Descent for Perceptrons

□ Input: $\mathcal{D} = \{\mathbf{x}^t, \mathbf{r}^t\}_{t=1}^N$

1. Initialize \mathbf{W} with small random numbers.
2. **repeat**
3. **for** each $(\mathbf{x}^t, \mathbf{r}^t) \in \mathcal{D}$ in random order
4. **for** each output node index i
5. $s_i \leftarrow \mathbf{w}_i^\top \mathbf{x}^t$
6. **end**
7. **for** each output node index i
8. $y_i \leftarrow \exp(s_i) / \sum_i \exp(s_i)$
9. **end**
10. **for** each output node index i
11. **for** each input node index j
12. $w_{ij} \leftarrow w_{ij} + \eta(r_i^t - y_i)x_j^t$
13. **end**
14. **end**
15. **end**
16. **until** convergence


□ Output: \mathbf{W}

; epoch

; regression: $y_i \leftarrow s_i$

sigmoid: $y \leftarrow 1/(1 + \exp(-\mathbf{w}^\top \mathbf{x}))$

Hard threshold: $y \leftarrow \begin{cases} 1 & \text{if } \mathbf{w}^\top \mathbf{x} > 0 \\ 0 & \text{otherwise} \end{cases}$



Contents


- Perceptrons
- Multilayer Perceptrons
 - Boolean Functions
 - Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications

Boolean AND

□ Boolean AND operation

- $y = x_1 \text{ AND } x_2$
- $y = \begin{cases} 1 & \text{if } \mathbf{w}^\top \mathbf{x} > 0 \\ 0 & \text{otherwise} \end{cases}$
- $\mathbf{w}^\top \mathbf{x} = w_0 + w_1 x_1 + w_2 x_2 > 0$
 $\Rightarrow x_2 > -\frac{w_1}{w_2} x_1 - \frac{w_0}{w_2}$

x_1	x_2	r
0	0	0
0	1	0
1	0	0
1	1	1




Boolean XOR

□ Boolean XOR operation

- $y = x_1 \text{ XOR } x_2$
- $y = \begin{cases} 1 & \text{if } \mathbf{w}^\top \mathbf{x} > 0 \\ 0 & \text{otherwise} \end{cases}$
- $\mathbf{w}^\top \mathbf{x} = w_0 + w_1 x_1 + w_2 x_2 > 0$
 $\Rightarrow x_2 > -\frac{w_1}{w_2} x_1 - \frac{w_0}{w_2}$


x_1	x_2	r
0	0	0
0	1	1
1	0	1
1	1	0



Boolean XOR

- Boolean XOR function as a disjunction of conjunctions
 - $x_1 \text{ XOR } x_2 = (x_1 \text{ AND } \neg x_2) \text{ OR } (\neg x_1 \text{ AND } x_2)$

x_1	x_2	z_1	z_2	r
0	0	0	0	0
0	1	0	1	1
1	0	1	0	1
1	1	0	0	0



Contents

- Perceptrons
- Multilayer Perceptrons
 - Boolean Functions
 - Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications


Multilayer Perceptrons

□ Multilayer perceptrons (MLP)

- Input x_j
- Hidden node $z_i = \frac{1}{1+\exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$; nonlinear activation function
- Weight $w_{z,ij}$ (between input x_j and hidden node z_i)
- Output $y_i = \begin{cases} \mathbf{w}_{o,i}^\top \mathbf{z} & ; z_0 \equiv 1 \\ \frac{\exp(s_i)}{\sum_i \exp(s_i)} & ; s_i = \mathbf{w}_{o,i}^\top \mathbf{z} \end{cases}$
- Weight $w_{o,ij}$ (between hidden node z_j and output node y_i)


□ Importance of non-linearity in hidden layers

- $\mathbf{z} = \mathbf{Wx}$
- $\mathbf{y} = \mathbf{Vz}$
- $\mathbf{y} = \mathbf{VWx} = \mathbf{Ux}$













Universal Approximation Theorem

- Any arbitrary function with continuous input and output can be approximated arbitrarily closely by a two-layer network given sufficient number of hidden units, proper nonlinearities, and weights.
- Proof
 - Fourier theorem



Decision Regions of MLPs

Structure	Types of Decision Regions	Exclusive-or Problem	Classes with Meshed Regions	Region Shapes
One Layer 	Half-Plane			
Two Layers 	Typically Convex			
Three Layers 	Arbitrary			



Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
 - Regression
 - Classification
 - Deep Neural Networks
- Applications

Error Backpropagation

□ Gradient descent


- $E(\theta | \mathbf{x}, \mathbf{r}) \equiv \frac{1}{2} \sum_i (r_i - y_i)^2$
- $E(\theta | \mathbf{x}, \mathbf{r}) \equiv - \sum_i r_i \log y_i$
- $\theta \leftarrow \theta - \eta \nabla E$



Learning for Multiple Linear Output Nodes

- Weight update for multiple linear output nodes


- $E(\mathbf{W}_z, \mathbf{W}_o | \mathbf{x}, \mathbf{r}) \equiv \sum_k \underbrace{\frac{1}{2}(r_k - y_k)^2}_{E_k}$; $y_k = \mathbf{w}_{o,k}^\top \mathbf{z}$, $z_i = \frac{1}{1 + \exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$
- $\frac{\partial E}{\partial w_{o,ij}} = \sum_k \frac{\partial E_k}{\partial y_i} \frac{\partial y_i}{\partial w_{o,ij}}$
 $= \underbrace{-(r_i - y_i)}_{\frac{\partial E_i}{\partial y_i}} \frac{z_j}{\frac{\partial y_i}{\partial w_{o,ij}}}$
- $w_{o,ij} \leftarrow w_{o,ij} + \eta(r_i - y_i)z_j$



Learning for Hidden Nodes

- Weight update for a hidden node connected to multiple linear output nodes

- $E(\mathbf{W}_z, \mathbf{W}_o | \mathbf{x}, \mathbf{r}) \equiv \sum_k \underbrace{\frac{1}{2}(r_k - y_k)^2}_{E_k}$; $y_k = \mathbf{w}_{o,k}^\top \mathbf{z}$, $z_i = \frac{1}{1 + \exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$
- $$\begin{aligned} \frac{\partial E}{\partial w_{z,ij}} &= \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_i} \frac{\partial z_i}{\partial w_{z,ij}} \\ &= \sum_k \underbrace{(r_k - y_k)}_{\frac{\partial E_k}{\partial y_k}} \underbrace{w_{o,ki}}_{\frac{\partial y_k}{\partial z_i}} \underbrace{z_i(1 - z_i)x_j}_{\frac{\partial z_i}{\partial w_{z,ij}}} \end{aligned}$$
- $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k - y_k) w_{o,ki} z_i (1 - z_i) x_j$




Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
 - Regression
 - Classification
 - Deep Neural Networks
- Applications

Learning for a Sigmoid Output Node

- Weight update for a single sigmoid output node


- $E(\mathbf{W}_z, \mathbf{w}_o | \mathbf{x}, r) \equiv -r \log y - (1 - r) \log(1 - y)$; $y = \frac{1}{1 + \exp(-s)}$, $s = \mathbf{w}_o^\top \mathbf{z}$
- $$\begin{aligned} \frac{\partial E}{\partial w_{o,i}} &= \frac{\partial E}{\partial y} \frac{\partial y}{\partial s} \frac{\partial s}{\partial w_{o,i}} \\ &= \underbrace{\left(-\frac{r}{y} + \frac{1-r}{1-y} \right)}_{\frac{\partial E}{\partial y}} \underbrace{y(1-y)}_{\frac{\partial y}{\partial s}} \underbrace{\frac{z_i}{\partial s}}_{\frac{\partial s}{\partial w_{o,i}}} \\ &= -(r - y) z_i \\ \mathbf{w}_{o,i} &\leftarrow \mathbf{w}_{o,i} + \eta(r - y) z_i \end{aligned}$$



Learning for Hidden Nodes

- Weight update for a hidden node connected to a single sigmoid output node


- $E(\mathbf{W}_z, \mathbf{w}_o | \mathbf{x}, r) \equiv -r \log y - (1 - r) \log(1 - y)$; $y = \frac{1}{1 + \exp(-\mathbf{w}_o^\top \mathbf{z})}$
- $$\begin{aligned} \frac{\partial E}{\partial w_{z,ij}} &= \frac{\partial E}{\partial y} \frac{\partial y}{\partial z_i} \frac{\partial z_i}{\partial w_{z,ij}} \\ &= \underbrace{\left(-\frac{r}{y} + \frac{1-r}{1-y} \right)}_{\frac{\partial E}{\partial y}} \underbrace{y(1-y)w_{o,i}}_{\frac{\partial y}{\partial z_i}} \underbrace{z_i(1-z_i)x_j}_{\frac{\partial z_i}{\partial w_{z,ij}}} \\ &= -(r - y)w_{o,i}z_i(1 - z_i)x_j \\ \mathbf{w}_{z,ij} &\leftarrow \mathbf{w}_{z,ij} + \eta(r - y)w_{o,i}z_i(1 - z_i)x_j \end{aligned}$$



Learning for Softmax Output Nodes

□ Weight update for a softmax output node


- $E(\mathbf{W}_z, \mathbf{W}_o | \mathbf{x}, \mathbf{r}) \equiv -\sum_k r_k \underbrace{\log y_k}_{E_k}$; $y_k = \frac{\exp(s_k)}{\sum_{\hat{k}} \exp(s_{\hat{k}})}$, $s_i = \mathbf{w}_{o,i}^\top \mathbf{z}$
- $\frac{\partial E}{\partial w_{o,ij}} = -\sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial s_i} \frac{\partial s_i}{\partial w_{o,ij}}$; $z_i = \frac{1}{1 + \exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$
- $= - \left(\sum_{k \neq i} \frac{r_k}{y_k} \underbrace{\left(-\frac{\exp(s_k) \exp(s_i)}{\left(\sum_{\hat{k}} \exp(s_{\hat{k}}) \right)^2} \right)}_{\frac{\partial E_k}{\partial y_k}} + \frac{r_i}{y_i} \underbrace{\left(\frac{\exp(s_i)}{\sum_{\hat{k}} \exp(s_{\hat{k}})} - \frac{\exp(s_i) \exp(s_i)}{\left(\sum_{\hat{k}} \exp(s_{\hat{k}}) \right)^2} \right)}_{\frac{\partial y_i}{\partial s_i}} \right) \frac{z_j}{\frac{\partial s_i}{\partial w_{o,ij}}}$
- $= - \left(\sum_{k \neq i} \frac{r_k}{y_k} (-y_k y_i) + \frac{r_i}{y_i} (y_i - y_i^2) \right) z_j$
- $= -(\sum_{k \neq i} r_k (-y_i) + r_i (1 - y_i)) z_j$
- $= -(\sum_k r_k (-y_i) + r_i) z_j$
- $= -(r_i - y_i) z_j$
- $w_{o,ij} \leftarrow w_{o,ij} + \eta(r_i - y_i) z_j$



Learning for Hidden Nodes

- Weight update for a hidden node connected to multiple softmax output nodes

- $E(\mathbf{W}_z, \mathbf{W}_o | \mathbf{x}, \mathbf{r}) \equiv -\sum_k r_k \underbrace{\log y_k}_{E_k}$; $y_k = \frac{\exp(s_k)}{\sum_{\hat{k}} \exp(s_{\hat{k}})}$, $s_i = \mathbf{w}_{o,i}^\top \mathbf{z}$
- $\frac{\partial E}{\partial w_{z,ij}} = -\sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_i} \frac{\partial z_i}{\partial w_{z,ij}}$; $z_i = \frac{1}{1 + \exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$
- $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k - y_k) w_{o,ki} z_i (1 - z_i) x_j$



Learning for Output Nodes

□ Weight update for a single sigmoid output node


- $w_{o,i} \leftarrow w_{o,i} + \eta(r - y)z_i$; $y = \frac{1}{1+\exp(-\mathbf{w}_{o,i}^\top \mathbf{z})}$, $z_i = \frac{1}{1+\exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$

□ Weight update for multiple softmax output nodes

- $w_{o,ij} \leftarrow w_{o,ij} + \eta(r_i - y_i)z_j$; $y_i = \frac{\exp(s_i)}{\sum_i \exp(s_i)}$, $s_i = \mathbf{w}_{o,i}^\top \mathbf{z}$


□ Weight update for multiple linear output nodes

- $w_{o,ij} \leftarrow w_{o,ij} + \eta(r_i - y_i)z_j$; $y_i = \mathbf{w}_{o,i}^\top \mathbf{z}$



Learning for Hidden Nodes


- Weight update for a hidden node connected to a single sigmoid output node
 - $w_{z,ij} \leftarrow w_{z,ij} + \eta(r - y)w_{o,i}z_i(1 - z_i)x_j$
- Weight update for a hidden node connected to multiple softmax output nodes
 - $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k - y_k)w_{o,ki} z_i(1 - z_i)x_j$
- Weight update for a hidden node connected to multiple linear output nodes
 - $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k - y_k)w_{o,ki} z_i(1 - z_i)x_j$



Stochastic Gradient Descent for MLPs

- Input: $\mathcal{D} = \{\mathbf{x}^t, \mathbf{r}^t\}_{t=1}^N$
 1. Initialize \mathbf{W}_z and \mathbf{W}_o with small random numbers.
 2. **repeat**
 3. **for** each $(\mathbf{x}^t, \mathbf{r}^t) \in \mathcal{D}$ in random order
 4. **for** each hidden node index i
 5. $z_i \leftarrow 1/(1 + \exp(-\mathbf{w}_{z,i}^\top \mathbf{x}^t))$
 6. **end**
 7. **for** each output node index i
 8. $s_i \leftarrow \mathbf{w}_{o,i}^\top \mathbf{z}$
 9. **end**
 10. **for** each output node index i
 11. $y_i \leftarrow \exp(s_i)/\sum_i \exp(s_i)$
 12. **end**
 13. **for** each hidden node index i
 14. **for** each input node index j
 15. $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k^t - y_k) w_{o,ki} z_i (1 - z_i) x_j^t$
 16. **end**
 17. **end**
 18. **for** each output node index i
 19. **for** each hidden node index j
 20. $w_{o,ij} \leftarrow w_{o,ij} + \eta (r_i^t - y_i) z_j$
 21. **end**
 22. **end**
 23. **end**
 24. **until** convergence
- Output: \mathbf{W}_z and \mathbf{W}_o

; sigmoid: $y \leftarrow 1/(1 + \exp(-\mathbf{w}^\top \mathbf{x}))$




Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
 - Regression
 - Classification
 - Deep Neural Networks
- Applications

Deep Neural Networks


- Deep neural network (DNN)
 - Multiple hidden layers
 - “deep and narrow” vs. “shallow and wide”
 - Representation learning ; more abstract concept
 - e.g., Image recognition: pixels → edges → corners → digits
- Deep learning



Learning for Multiple Hidden Layers

- Weight update for a hidden node of the first hidden layer in a two hidden layer neural network with multiple linear output nodes


- $E(\mathbf{W}_1, \mathbf{W}_2, \mathbf{W}_3 | \mathbf{x}, \mathbf{r}) \equiv \sum_k \underbrace{\frac{1}{2}(r_k - y_k)^2}_{E_k}$; $y_k = \sum_m w_{3,km} z_{2,m}$
- $\frac{\partial E}{\partial w_{1,ij}} = \sum_k \frac{\partial E_k}{\partial y_k} \underbrace{\frac{\partial y_k}{\partial z_{1,i}} \frac{\partial z_{1,i}}{\partial w_{1,ij}}}_{\frac{\partial \sum_m w_{3,km} z_{2,m}}{\partial z_{1,i}}} = \sum_m \frac{\partial w_{3,km} z_{2,m}}{\partial z_{2,m}} \frac{\partial z_{2,m}}{\partial z_{1,i}}$; $z_{2,m} = \frac{1}{1 + \exp(-\sum_i w_{2,mi} z_{1,i})}$
 $= \sum_k \underbrace{(r_k - y_k)}_{\frac{\partial E_k}{\partial y_k}} \sum_m \underbrace{w_{3,km}}_{\frac{\partial w_{3,km} z_{2,m}}{\partial z_{2,m}}} \underbrace{z_{2,m}(1 - z_{2,m})}_{\frac{\partial z_{2,m}}{\partial z_{1,i}}} \underbrace{w_{2,mi}}_{\frac{\partial z_{1,i}}{\partial w_{1,ij}}} \underbrace{z_{1,i}(1 - z_{1,i})x_j}_{\frac{\partial z_{1,i}}{\partial w_{1,ij}}}$; $z_{1,i} = \frac{1}{1 + \exp(-\sum_j w_{1,ij} x_j)}$
- $w_{1,ij} \leftarrow w_{1,ij} + \eta \sum_k (r_k - y_k) \sum_m w_{3,km} z_{2,m} (1 - z_{2,m}) w_{2,mi} z_{1,i} (1 - z_{1,i}) x_j$



Learning for Multiple Hidden Layers

- Weight update for a hidden node of the first hidden layer in a two hidden layer neural network with multiple softmax output nodes
 - $E(\mathbf{W}_1, \mathbf{W}_2, \mathbf{W}_3 | \mathbf{x}, \mathbf{r}) \equiv - \sum_k \underbrace{r_k \log y_k}_{E_k} ; y_k = \frac{\exp(s_k)}{\sum_{\hat{k}} \exp(s_{\hat{k}})}, s_k = \sum_m w_{3,km} z_{2,m}$
 - $\frac{\partial E}{\partial w_{1,ij}} = - \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_{1,i}} \frac{\partial z_{1,i}}{\partial w_{1,ij}}$
 $; z_{2,m} = \frac{1}{1 + \exp(- \sum_i w_{2,mi} z_{1,i})}$
 - $w_{1,ij} \leftarrow w_{1,ij} + \eta \sum_k (r_k - y_k) \sum_m w_{3,km} z_{2,m} (1 - z_{2,m}) w_{2,mi} z_{1,i} (1 - z_{1,i}) x_j$


- Time complexity
 - $\mathcal{O}(n^L)$



Learning for DNNs

□ Weight update for an output node

- $\frac{\partial E}{\partial w_{3,ij}} = \sum_k \frac{\partial E_k}{\partial y_i} \frac{\partial y_i}{\partial w_{3,ij}} = \sum_k \frac{\partial E_k}{\partial y_i} \frac{\partial y_i}{\partial s_{3,i}} \frac{\partial s_{3,i}}{\partial w_{3,ij}}$
- $w_{3,ij} \leftarrow w_{3,ij} + \eta \underbrace{(r_i - y_i)}_{; s_{l,i} \equiv \sum_j w_{l,ij} z_{l-1,j}} z_{2,j}$
- $w_{3,ij} \leftarrow w_{3,ij} + \eta \delta_{3,i} z_{2,j} ; \delta_{3,i} \equiv r_i - y_i = \frac{\partial E}{\partial s_{3,i}}$



□ Weight update for a hidden node connected to multiple output nodes

- $\frac{\partial E}{\partial w_{2,ij}} = \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_{2,i}} \frac{\partial z_{2,i}}{\partial w_{2,ij}} = \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_{2,i}} \frac{\partial z_{2,i}}{\partial s_{2,i}} \frac{\partial s_{2,i}}{\partial w_{2,ij}}$
- $w_{2,ij} \leftarrow w_{2,ij} + \eta \underbrace{\sum_k (r_k - y_k) w_{3,ki} z'_{2,i}}_{; \delta_{2,i} \equiv \sum_k \delta_{3,k} w_{3,ki} z'_{2,i}} z_{1,j}$
- $w_{2,ij} \leftarrow w_{2,ij} + \eta \delta_{2,i} z_{1,j} ; \delta_{2,i} \equiv \sum_k \delta_{3,k} w_{3,ki} z'_{2,i} = \frac{\partial E}{\partial s_{2,i}}$

□ Weight update for a hidden node of the first hidden layer in a two hidden layer neural network with multiple output nodes

- $\frac{\partial E}{\partial w_{1,ij}} = \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_{1,i}} \frac{\partial z_{1,i}}{\partial w_{1,ij}} = \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_{1,i}} \frac{\partial z_{1,i}}{\partial s_{1,i}} \frac{\partial s_{1,i}}{\partial w_{1,ij}}$
- $w_{1,ij} \leftarrow w_{1,ij} + \eta \underbrace{\sum_k (r_k - y_k) \sum_m w_{3,km} z'_{2,m} w_{2,mi} z'_{1,i} x_j}_{; \delta_{1,i} \equiv \sum_m \delta_{2,m} w_{2,mi} z'_{1,i}} ; \delta_{1,i} \equiv \sum_m \delta_{2,m} w_{2,mi} z'_{1,i} = \frac{\partial E}{\partial s_{1,i}}$
- $w_{1,ij} \leftarrow w_{1,ij} + \eta \underbrace{\sum_m \sum_k (r_k - y_k) w_{3,km} z'_{2,m} w_{2,mi} z'_{1,i} x_j}_{; \delta_{1,i} \equiv \sum_m \delta_{2,m} w_{2,mi} z'_{1,i}} ; \delta_{1,i} \equiv \sum_m \delta_{2,m} w_{2,mi} z'_{1,i} = \frac{\partial E}{\partial s_{1,i}}$



Learning for DNNs

- Weight update for an output node

- $w_{L,ij} \leftarrow w_{L,ij} + \eta \underbrace{\frac{\partial E}{\partial s_{l,i}}}_{\frac{\delta E}{\partial s_{l,i}}} \underbrace{\frac{\partial s_{l,i}}{\partial w_{L,ij}}}_{z_{L-1,j}}$; $\delta_{L,i} \equiv r_i - y_i$, $\partial s_{l,i} \equiv \sum_j w_{l,ij} z_{l-1,j}$

- Weight update for a hidden node of the l -th hidden layer in a multiple hidden layer neural network

- $w_{l,ij} \leftarrow w_{l,ij} + \eta \underbrace{\frac{\partial E}{\partial s_{l,i}}}_{\frac{\delta E}{\partial s_{l,i}}} \underbrace{\frac{\partial s_{l,i}}{\partial w_{l,ij}}}_{z_{l-1,j}}$; $\delta_{l,i} \equiv \sum_m \delta_{l+1,m} w_{l+1,mi} z'_{l,i}$





Stochastic Gradient Descent for DNNs

□ Input: $\mathcal{D} = \{\mathbf{x}^t, \mathbf{r}^t\}_{t=1}^N$

```

1. Initialize  $\{\mathbf{W}_l\}_{l=1}^L$  with small random numbers.
2. repeat
3.   for each  $(\mathbf{x}^t, \mathbf{r}^t) \in \mathcal{D}$  in random order
4.     for each hidden layer index  $l$  from 1 to  $L - 1$  ; forward process: feed forward
5.       for each node index  $i$  in the  $l$ -th layer
6.          $z_{l,i} \leftarrow 1/(1 + \exp(-\mathbf{w}_{l,i}^\top \mathbf{z}_{l-1}))$  ;  $\mathbf{z}_0 \equiv \mathbf{x}^t$ 
7.       end
8.     end
9.     for each output node index  $i$ 
10.     $y_i \leftarrow \exp(\mathbf{w}_{L,i}^\top \mathbf{z}_{L-1}) / \sum_i \exp(\mathbf{w}_{L,i}^\top \mathbf{z}_{L-1})$  } ;  $y_i \leftarrow \mathbf{w}_{L,i}^\top \mathbf{z}_{L-1}$ 
11.    end ;  $y \leftarrow 1/(1 + \exp(-\mathbf{w}_L^\top \mathbf{z}_{L-1}))$ 
12.    for each output node index  $i$ 
13.       $\delta_{L,i} \leftarrow r_i^t - y_i$ 
14.    end
15.    for each layer index  $l$  from  $L$  to 1 ; backward process
16.      for each node index  $j$  in the  $(l - 1)$ -th layer
17.        if  $l \neq 1$  then  $\delta_{l-1,j} \leftarrow \sum_i \delta_{l,i} w_{l,ij} z'_{l-1,j}$  ; error back-propagation
18.        for each node index  $i$  in the  $l$ -th layer
19.           $w_{l,ij} \leftarrow w_{l,ij} + \eta \delta_{l,i} z_{l-1,j}$ 
20.        end
21.      end
22.    end
23.  end
24. until convergence

```





; $\mathbf{z}_0 \equiv \mathbf{x}^t$

; $y_i \leftarrow \mathbf{w}_{L,i}^\top \mathbf{z}_{L-1}$
; $y \leftarrow 1/(1 + \exp(-\mathbf{w}_L^\top \mathbf{z}_{L-1}))$

; backward process

; error back-propagation




□ Output: all $\{\mathbf{W}_l\}_{l=1}^L$

Contents


- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications
 - Multilabel Classification
 - Dimensionality Reduction
 - Representation Learning

Multilabel Classification

- K independent two-class classification



- Multitask learning



Learning for Multilabel Classification

□ K dependent two-class classification

- $E(\mathbf{W}_z, \mathbf{W}_o | \mathbf{x}, \mathbf{r}) \equiv \sum_k \underbrace{[-r_k \log y_k - (1 - r_k) \log(1 - y_k)]}_{E_k}$

- Weight update for K sigmoid output nodes

- $$\begin{aligned} \frac{\partial E}{\partial w_{o,ij}} &= \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial w_{o,ij}} \\ &= \underbrace{\left(-\frac{r_i}{y_i} + \frac{1-r_i}{1-y_i} \right)}_{\frac{\partial E_i}{\partial y_i}} \underbrace{y_i(1-y_i)z_j}_{\frac{\partial y_i}{\partial w_{o,ij}}} \\ &= -(r_i - y_i)z_j \end{aligned}$$

- $w_{o,ij} \leftarrow w_{o,ij} + \eta(r_i - y_i)z_j$

- Weight update for a hidden node connected to K sigmoid output nodes


- $$\begin{aligned} \frac{\partial E}{\partial w_{z,ij}} &= \sum_k \frac{\partial E_k}{\partial y_k} \frac{\partial y_k}{\partial z_i} \frac{\partial z_i}{\partial w_{z,ij}} \\ &= \sum_k \underbrace{\left(-\frac{r_k}{y_k} + \frac{1-r_k}{1-y_k} \right)}_{\frac{\partial E_k}{\partial y_k}} \underbrace{y_k(1-y_k)w_{o,ki}}_{\frac{\partial y_k}{\partial z_i}} \underbrace{z_i(1-z_i)x_j}_{\frac{\partial z_i}{\partial w_{z,ij}}} \\ &= -\sum_k (r_k - y_k)w_{o,ki} z_i(1-z_i)x_j \end{aligned}$$

- $w_{z,ij} \leftarrow w_{z,ij} + \eta \sum_k (r_k - y_k)w_{o,ki} z_i(1-z_i)x_j$

$$r_k = \begin{cases} 1 & \text{if } x \text{ has label } k \\ 0 & \text{otherwise} \end{cases}$$

$$y_k = \frac{1}{1+\exp(-\mathbf{w}_{o,k}^\top \mathbf{z})}$$

$$z_i = \frac{1}{1+\exp(-\mathbf{w}_{z,i}^\top \mathbf{x})}$$



Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications
 - Multilabel Classification
 - Dimensionality Reduction
 - Representation Learning


Autoencoders

□ Autoencoders [Cottrell, Munro, and Zipser 1987]

- $\mathbf{z} = \text{ENCODE}(\mathbf{x} | \mathbf{W})$
- $\tilde{\mathbf{x}} = \text{DECODE}(\mathbf{z} | \mathbf{V})$


□ Reconstruction error of an autoencoder

- $$\begin{aligned} E(\mathbf{W}, \mathbf{V} | \mathcal{D}) &= \frac{1}{N} \sum_t \|\mathbf{x}^t - \tilde{\mathbf{x}}^t\|^2 \quad ; \|\mathbf{x}\|_2 \equiv \sqrt{\sum_i x_i^2} \text{ (Euclidean norm or } \ell^2 \text{ norm)} \\ &= \frac{1}{N} \sum_t \|\mathbf{x}^t - \text{DECODE}(\text{ENCODE}(\mathbf{x}^t | \mathbf{W}) | \mathbf{V})\|^2 \end{aligned}$$



Dimensionality Reduction Using Autoencoders

- Dimensionality reduction using linear autoencoders




Dimensionality Reduction Example




Variants of Autoencoders

- Denoising autoencoder [Vincent *et al.* 2008]
 - Noiseless and noisy versions of the same instance with the same desired output
- Sparse autoencoder [Ranzato *et al.* 2007]
 - dimensionality of input < dimensionality of hidden layer
- Deep autoencoder and stacked autoencoder




Dimensionality Reduction Example



Multidimensional Scaling

- Multidimensional scaling (MDS)
 - high-dimensional space \Rightarrow low-dimensional space
 - Trying to keep the inter-point distances unchanged
 - Dimensionality reduction
- Sammon mapping
 - $\mathbf{z} = \mathbf{g}(\mathbf{x}|\theta)$ where $\mathbf{x} \in \mathcal{R}^d$, $\mathbf{z} \in \mathcal{R}^k$, and $k < d$
 - Sammon stress: $\sum_{t,s} \frac{(\|\mathbf{g}(\mathbf{x}^t|\theta) - \mathbf{g}(\mathbf{x}^s|\theta)\| - \|\mathbf{x}^t - \mathbf{x}^s\|)^2}{\|\mathbf{x}^t - \mathbf{x}^s\|^2} = \sum_{t,s} \frac{(\|\mathbf{z}^t - \mathbf{z}^s\| - \|\mathbf{x}^t - \mathbf{x}^s\|)^2}{\|\mathbf{x}^t - \mathbf{x}^s\|^2}$
- An MLP with d inputs, hidden units, and $k < d$ output units



Contents

- Perceptrons
- Multilayer Perceptrons
- Error Backpropagation Algorithm
- Applications
 - Multilabel Classification
 - Dimensionality Reduction
 - Representation Learning


Representation Learning

□ Linear model and nonlinear model

- $y = \sum_i v_i x_i + v_0$
- $y = \sum_i v_i \phi_i(\mathbf{x})$; ϕ_i : basis function
e.g., polynomial basis function

□ Embedding (also called code or *representation*)


- $y = \sum_i v_i \phi(\mathbf{x}|\mathbf{w}_i)$; $\phi(\mathbf{x}|\mathbf{w}_i) \equiv \frac{1}{1+\exp(-\mathbf{w}_i^\top \mathbf{x})}$
- Learnable basis function
- A new space where the problem becomes linear, e.g., XOR
- Unsupervised learning, e.g., autoencoder




- Dimensionality reduction: dimensionality of embedding < dimensionality of input
- Transfer learning: using the embedding trained on one task as input to another task
- Multitask learning: learning similar tasks with shared hidden layers
- Self-supervised learning and finetuning: using the embedding trained on a large unlabeled data set as the input in supervised learning with a small labeled data set

Word Embedding

- Binary vector representation of words
 - d -dimensional one-hot vector ; bag of words
 - High dimensional
 - No syntactic/semantic relationship
- Continuous vector representation of words (WORD2VEC, Mikolov *et al.* 2013)
 - Input: context (windows of words)
 - Each word in the window at a time (SKIP-GRAM)
e.g., “Protesters in Paris clash with the French police.”
 - Hidden layer: less than d hidden units
 - Output: center word (d -dimensional one-hot vector)



Word Embedding Example



$$\text{VEC}(\text{"French"}) - \text{VEC}(\text{"Paris"}) + \text{VEC}(\text{"London"}) = \text{VEC}(\text{"English"})$$

Summary and Preview

- Perceptrons
 - Introduction
 - Learning for Perceptrons
- Multilayer Perceptrons
 - Boolean Functions
 - Multilayer Perceptrons
- Error Backpropagation Algorithm
 - Regression
 - Classification
 - Deep Neural Networks
- Applications
 - Multilabel Classification
 - Dimensionality Reduction
 - Representation Learning
- Architecture of DNNs